



INFO 7250 – ENGINEERING OF BIG DATA SYSTEMS SEC 01

ANALYSING IOWA LIQUOR SALES USING HADOOP

GAURESH CHAVAN
001854530

INDEX

1. Introduction

- 1.1. Problem statement
- 1.2. Data glossary

2. Technology

- 3.1. Hadoop
- 3.2. Pig
- 3.3. Hive

3. Use cases

- 3.1. Analyzing liquor sales during Holiday season vs Regular season
- 3.2. Assessing top 5 liquor labels based on their sales from 2012 through 2019
- 3.3. Investigating alcohol category performance over the years
- 3.4. Exploring expansion opportunities through zip code analysis
- 3.5. Researching the top alcohol brand sold by each store in the state

INTRODUCTION

Problem Statement:

The objective of this project is to explore the dataset and implement Hadoop ecosystem components on it. We are primarily going to utilize Hadoop's MapReduce , Pig and Hive to derive some analytic insights from data.

Data Glossary:

This dataset contains the spirits purchase information of Iowa Class "E" liquor licenses by product and date of purchase from January 1, 2012 to current. The dataset can be used to analyze total spirits sales in Iowa of individual products at the store level.

Updated: November 2, 2019

Data provided by: Iowa Department of Commerce, Alcoholic Beverages Division

This dataset has 17.3 Million rows across 17 columns. Each row represents a transaction.

- Invoice/Item Number
- Date
- Store Number
- Store Name
- Address
- City
- Zip Code
- Store Location
- County Number
- County
- Category
- Category Name
- Vendor Number
- Vendor Name
- Item Number
- Item Description
- Pack
- Bottle Volume (ml)
- State Bottle Cost
- State Bottle Retail
- Bottles Sold
- Sale (Dollars)
- Volume Sold (Liters)
- Volume Sold (Gallons)

TECHNOLOGY



Apache Hive



Apache MapReduce:

MapReduce is a programming model and an associated implementation for processing and generating big data sets with a parallel, distributed algorithm on a cluster.

Code for MapReduce is written in Java. Eclipse IDE is used as a development environment.

Implemented on Use Case 3.1 & 3.2

Apache Pig:

Apache Pig is a high-level platform for creating programs that run on Apache Hadoop. The language for this platform is called Pig Latin. Pig can execute its Hadoop jobs in MapReduce, Apache Tez, or Apache Spark.

Implemented on Use Case 3.3 & 3.4

Apache Hive:

Apache Hive is a data warehouse software project built on top of Apache Hadoop for providing data query and analysis. Hive gives a SQL-like interface to query data stored in various databases and file systems that integrate with Hadoop.

Implemented on Use Case 3.5

Both Apache Pig & Hive is being used owing to its fast processing and slow learning curve. While both are languages (Pig is a data flow language, Hive provides SQL constructs), there is a little learning curve and reduces the number of lines of code mapper/reducer would otherwise take.

USE CASES

1. Analyzing liquor sales during Holiday Season vs Regular Season

It would not be surprising to assume, people naturally consume more alcohol during holiday season i.e. owing to Thanksgiving, Hanukkah, Christmas and New Year's. So, it would be interesting to see if our assumption is correct. We will be comparing sales for first 10 months i.e. January – October to those for last 2 months i.e. November - December for every year from 2012- 2019.

Our comparison metrics will be **Average Monthly Sales & Average sale per transaction** for that year

Mapper:

```
23     try {
24         if (key.get() == 0 && value.toString().contains("Item Description"))
25             return;
26         else
27         {
28             String[] tokens = line.split(",");
29
30             SimpleDateFormat sdf = new SimpleDateFormat("MM/dd/yyyy");
31             Date new_date = sdf.parse(tokens[1]); //parsing into Date Format
32
33             SimpleDateFormat sdf2 = new SimpleDateFormat("yyyy");
34             String tx_year = sdf2.format(new_date); //extracting year
35
36             SimpleDateFormat sdf3 = new SimpleDateFormat("MM");
37             String tx_month = sdf3.format(new_date); //extracting month
38
39             year.set(tx_year); //setting key
40
41             String tx_sales = tokens[21].replace("$", ""); //getting sales
42
43             if(tx_sales != null && !tx_sales.equals("")) {
44                 sales = Double.parseDouble(tx_sales); //setting value
45             }
46
47             String type = "R";
48
49             if(tx_month.equals("11") || tx_month.equals("12")) //checking to see holiday months: November or December
50             {
51                 type = "H";
52             }
53
54             cw = new CustomWritable(sales,type);
55
56             context.write(year, cw);
57         }
58     }
```

Since the date we have in our data is a string, we parse it into *Date()* format. Then we extract 'Year' which is our key, and 'month' which will be used to grab key-value pairs as 'Holiday' sales. Also, we define **a type == 'H', 'R'** which acts as a filter to separate the sales when we reduce them.

Reducer:

In reducer, we define two separate counts and sums; one of holiday and other for regular average. If type = H, we calculate monthly average as (sum/2) i.e. for November &

December, else we calculate regular monthly average as (sum/10) i.e. for January – October.

```
7 public class reducer extends Reducer<Text, CustomWritable, Text, Text>
8 {
9
10     public void reduce(Text key, Iterable<CustomWritable> values, Context context) throws IOException, InterruptedException
11     {
12         double sumH = 0;
13         double sumR = 0;
14
15         int countH = 0;
16         int countR = 0;
17
18         for (CustomWritable x : values)
19         {
20             if (x.type.equals("H"))
21             {
22                 sumH += x.sales;
23                 countH++;
24             }
25             else
26             {
27                 sumR += x.sales;
28                 countR++;
29             }
30         }
31
32         double holidayMonthlyAverage = (sumH/2);
33         double holidayTxAverage = (sumH/countH);
34
35         double regularMonthlyAverage = (sumR/10);
36         double regularTxAverage = (sumR/countR);
37
38         String out = "Hday_MM Avg: " + holidayMonthlyAverage + ", Hday_Tx Avg: "
39         + holidayTxAverage + ", Reg_MM Avg: " + regularMonthlyAverage + ", Reg_Tx Avg: " + regularTxAverage;
40
41         context.write(key, new Text(out));
42     }
}
```

Output:

```
2012 Hday_MM Avg: 1.8346200149999086E7, Hday_Tx Avg: 98.55653347586657, Reg_MM Avg: 1.626766919297742E7, Reg_Tx Avg: 95.14586654495815
2013 Hday_MM Avg: 1.710375582000114E7, Hday_Tx Avg: 94.56563719650207, Reg_MM Avg: 1.681613573796019E7, Reg_Tx Avg: 98.80046613726074
2014 Hday_MM Avg: 1.9551097180004068E7, Hday_Tx Avg: 100.5577798294166, Reg_MM Avg: 1.6687745410984691E7, Reg_Tx Avg: 97.64951441320567
2015 Hday_MM Avg: 2.0338168035003364E7, Hday_Tx Avg: 99.95757581046338, Reg_MM Avg: 1.7112676971995074E7, Reg_Tx Avg: 96.27130518627678
2016 Hday_MM Avg: 2.7070674755003843E7, Hday_Tx Avg: 128.4656860191049, Reg_MM Avg: 1.8591404729980486E7, Reg_Tx Avg: 100.03731465028858
2017 Hday_MM Avg: 2.6771955890000556E7, Hday_Tx Avg: 131.17784834067763, Reg_MM Avg: 2.341175087297933E7, Reg_Tx Avg: 124.32571683990601
2018 Hday_MM Avg: 2.9552565225002263E7, Hday_Tx Avg: 141.06000723139357, Reg_MM Avg: 2.4815640939981773E7, Reg_Tx Avg: 128.14349294173905
2019 Hday_MM Avg: 0.0, Hday_Tx Avg: NaN, Reg_MM Avg: 2.635629230098617E7, Reg_Tx Avg: 135.4323798856996
```

2. Assessing top 5 liquor labels based on their sales from 2012-2019

WritableComparable:

To calculate top 5 liquor brands per year by sales, we create a CustomWritable class where we define a composite key (year, brand) with value being (sales).

The compareTo() is WritableComparable's natural comparison method. It returns 0 if both objects are same and with a number higher or lower you determine the order between your objects. In our case, we compare our composite key elements: year, brand.

```

9 public class CustomWritable implements WritableComparable<Object> {
10
11     String year;
12     String brand;
13
14     public CustomWritable(String year, String brand)
15     {
16         this.year = year;
17         this.brand = brand;
18     }
19
20     public CustomWritable()
21     {}
22
23     public void readFields(DataInput in) throws IOException {
24         this.year = in.readUTF();
25         this.brand = in.readUTF();
26     }
27
28     public void write(DataOutput out) throws IOException {
29         out.writeUTF(this.year);
30         out.writeUTF(this.brand);
31     }
32
33     @Override
34     public String toString() {
35         return "Year: " + this.year + ", Brand: " + this.brand;
36     }
37
38     public int compareTo(Object wc2) {
39         CustomWritable key2 = (CustomWritable) wc2;
40         int intCount = this.year.compareTo(key2.year);
41         return intCount == 0 ? this.brand.compareTo(key2.brand) : intCount;
42     }
43
44     @Override
45     public int hashCode()
46     {
47         final int prime = 31;
48         int result = 1;
49         result = prime * result + ((this.year == null) ? 0 : this.year.hashCode());
50         result = prime * result + ((this.brand == null) ? 0 : this.brand.hashCode());
51         return result;
52     }

```

Partitioner:

By default, the partitioner implementation is called HashPartitioner. It uses the hashCode() method of the key objects modulo the number of partitions total to determine which partition to send a given (key, value) pair to. The getPartition() method receives a key and a value and the number of partitions to split the data, a number in the range [0, numPartitions) must be returned by this method, indicating which partition to send the key and value to.

```

6 public class YearPartitioner extends Partitioner<CustomWritable, DoubleWritable> {
7
8     @Override
9     public int getPartition(CustomWritable key, DoubleWritable value, int numPartitions)
10     {
11
12         return Math.abs(key.year.hashCode()) % numPartitions;
13
14         // return Integer.parseInt(key.year) % numPartitions;
15     }
16 }
17 }

```

In our reducer, we define two tree maps: one to store count of sales of all the brands sold in a given year; other to store the first map per year i.e. for first tree map, brand is the key whereas for second map, year will behave as a key.

Reducer:

```
14 public class reducer extends Reducer<CustomWritable, DoubleWritable, CustomWritable, DoubleWritable>
15 {
16     private TreeMap<String, Double> tmap2;
17     private TreeMap<String, TreeMap<String, Double>> yearMap;
18     String year;
19     @Override
20     public void setup(Context context) throws IOException, InterruptedException
21     {
22         tmap2 = new TreeMap<String, Double>();
23         yearMap = new TreeMap<String, TreeMap<String, Double>>();
24     }
25
26     public void reduce(CustomWritable key, Iterable<DoubleWritable> values, Context context) throws IOException, InterruptedException
27     {
28         double sum = 0;
29
30         String[] keySplits = key.toString().split(","); // 0: Year: <year>, 1: Liquor: <brand>
31         year = keySplits[0].split(":")[1].trim();
32         String brand = keySplits[1].split(":")[1].trim();
33
34         if(yearMap.get(year) == null)
35         {
36             tmap2 = new TreeMap<String, Double>();
37         }
38         else
39         {
40             tmap2 = yearMap.get(year);
41         }
42
43         if(tmap2.get(brand) == null)
44         {
45             tmap2.put(brand, sum);
46         }
47         else
48         {
49             sum = tmap2.get(brand);
50         }
51
52         for (DoubleWritable x : values)
53         {
54             sum += x.get();
55         }
56
57         tmap2.put(brand, sum);
58         yearMap.put(year, tmap2);
59     }
}
```

Cleanup:

The first step in clean up is to sort the values based on sales (value) coming from year map which is why we see it as an entry set to a Sorted Set. The **Java Sorted Set** interface behaves like a normal **Set** with the exception that the elements it contains are **sorted** internally. This means that when you iterate the elements of a **Sorted Set** the elements are iterated in the **sorted** order.

In a way sorted set assumes sorting thus relieving us of using a sort comparator. Second, we initialize a count = 0 which is then used to iterate over the sorted set to print out only 5 top values which is our desired result.


```

61@ Override
62 protected void cleanup(Reducer<CustomWritable, DoubleWritable, CustomWritable, DoubleWritable>.Context context) throws IOException, InterruptedException
63 {
64
65     for(Map.Entry<String, TreeMap<String, Double>> yearEntry: yearMap.entrySet())
66     {
67         int count = 0;
68         TreeMap<String, Double> yearValueMap = yearEntry.getValue();
69
70         SortedSet<Map.Entry<String, Double>> sortedset = new TreeSet<Map.Entry<String, Double>>
71         (new Comparator<Map.Entry<String, Double>>()
72         {
73             public int compare(Map.Entry<String, Double> e1, Map.Entry<String, Double> e2)
74             {
75                 return e2.getValue().compareTo(e1.getValue());
76             }
77         });
78
79         sortedset.addAll(yearValueMap.entrySet());
80
81         CustomWritable cw;
82
83         for(Map.Entry<String, Double> entry : sortedset)
84         {
85             if(count == 5)
86             {
87                 break;
88             }
89             count++;
90             cw = new CustomWritable(yearEntry.getKey(), entry.getKey());
91             context.write(cw, new DoubleWritable(entry.getValue()));
92         }
93         context.write(null, null);
94     }
95 }
96
97 }
98
99

```

Output:

```

Year: 2012, Brand: LUXCO-ST LOUIS      1.8151307680001702E7
Year: 2012, Brand: BROWN-FORMAN CORPORATION  1.5119238469999464E7
Year: 2012, Brand: PROXIMO      9307097.990001196
Year: 2012, Brand: SIDNEY FRANK IMPORTING CO.  5222647.730000268
Year: 2012, Brand: WILSON DANIELS LTD.  5197392.619999735
Year: 2013, Brand: LUXCO-ST LOUIS      1.874094039999595E7
Year: 2013, Brand: BROWN-FORMAN CORPORATION  1.6090561410001062E7
Year: 2013, Brand: PROXIMO      8115770.149998465
Year: 2013, Brand: SIDNEY FRANK IMPORTING CO.  4822736.600000324
Year: 2013, Brand: WILSON DANIELS LTD.  2820783.6099999268
Year: 2014, Brand: LUXCO-ST LOUIS      1.7902289779997893E7
Year: 2014, Brand: BROWN-FORMAN CORPORATION  1.678282342999734E7
Year: 2014, Brand: PROXIMO      7821755.619998185
Year: 2014, Brand: SIDNEY FRANK IMPORTING CO.  4582258.740000276
Year: 2014, Brand: CAMPARI(SKYY)  2725793.2099999716
Year: 2015, Brand: BROWN-FORMAN CORPORATION  1.80513852899985E7
Year: 2015, Brand: LUXCO-ST LOUIS      1.6660449889997894E7
Year: 2015, Brand: PROXIMO      8279668.309998219
Year: 2015, Brand: SIDNEY FRANK IMPORTING CO.  4023823.6100003105
Year: 2015, Brand: FIFTH GENERATION INC.  3766131.0199997365
Year: 2016, Brand: LUXCO-ST LOUIS      1.0713402419997789E7
Year: 2016, Brand: BROWN-FORMAN CORPORATION  9256091.699999588
Year: 2016, Brand: PROXIMO      8970297.479999341
Year: 2016, Brand: BROWN FORMAN CORP.  6382450.350000258
Year: 2016, Brand: FIFTH GENERATION INC.  3780106.0399998245
Year: 2017, Brand: BROWN FORMAN CORP.  1.5246581549998695E7
Year: 2017, Brand: PROXIMO      9608637.38000276
Year: 2017, Brand: E & J GALLO WINERY  6087477.640000242
Year: 2017, Brand: LAIRD & COMPANY  4574345.310000283
Year: 2017, Brand: MCCORMICK DISTILLING CO.  3152896.0599994203
Year: 2018, Brand: BROWN FORMAN CORP.  1.6127212409998564E7
Year: 2018, Brand: PROXIMO      9956119.180003
Year: 2018, Brand: E & J GALLO WINERY  5968684.570000038
Year: 2018, Brand: LAIRD & COMPANY  4369411.010000804
Year: 2018, Brand: MCCORMICK DISTILLING CO.  3563372.82999919
Year: 2019, Brand: BROWN FORMAN CORP.  1.325229781999915E7
Year: 2019, Brand: PROXIMO      9286240.540002625
Year: 2019, Brand: E & J GALLO WINERY  4749800.520000029
Year: 2019, Brand: LAIRD & COMPANY  3442661.1800007275
Year: 2019, Brand: MCCORMICK DISTILLING CO.  2935071.189999524

```

3. Investigating alcohol category performance over the years

Pig Script:

```
sales = LOAD 'hdfs://localhost:9000/project_data/Iowa_Liquor_Sales.csv'
        using org.apache.pig.piggybank.storage.CSVExcelStorage(',', 'NO_MULTILINE',
        'NOCHANGE', 'SKIP_INPUT_HEADER')
        as (Invoice:chararray, Date:chararray, Store_name:chararray, Store_number:int,
        Address:chararray, City:chararray, Zipcode:int, Store_location:chararray,
        County_Number:int, County:chararray, Category:int, Category_name:chararray,
        Vendor_number:int, Vendor_name:chararray, Item_number:int, Item_description:chararray,
        Pack:int, Bottle_volume:int, State_bottle_cost:float, State_bottle_retail:float,
        Bottles_sold:int, Sale:float, Volume_liters:float, Volume_gallons:float);

category = GROUP sales by UPPER(Category_name);

op = FOREACH category GENERATE group as category,
                                SUM(sales.Sale) as total_sales,
                                SUM(sales.Bottles_sold) as bottles_sold,
                                SUM(sales.Volume_liters) as total_volume;

STORE op INTO 'hdfs://localhost:9000/pig_outputs/categoryAnalysis_op';
```

This is a straightforward script.

1. Load data into 'sales' (while skipping the header row)
2. Group by Category name
3. For each category name, compute total sales, total bottles sold, and total volume sold
4. Store output to HDFS

Pig Runtime:

```
2019-12-08 06:57:50,875 [main] INFO org.apache.hadoop.yarn.client.RMProxy - Connecting to ResourceManager at /0.0.0.0:8032
2019-12-08 06:57:50,878 [main] INFO org.apache.hadoop.mapred.ClientServiceDelegate - Application state is completed. FinalApplicationStat
2019-12-08 06:57:50,924 [main] INFO org.apache.hadoop.yarn.client.RMProxy - Connecting to ResourceManager at /0.0.0.0:8032
2019-12-08 06:57:50,927 [main] INFO org.apache.hadoop.mapred.ClientServiceDelegate - Application state is completed. FinalApplicationStat
2019-12-08 06:57:50,950 [main] INFO org.apache.hadoop.yarn.client.RMProxy - Connecting to ResourceManager at /0.0.0.0:8032
2019-12-08 06:57:50,954 [main] INFO org.apache.hadoop.mapred.ClientServiceDelegate - Application state is completed. FinalApplicationStat
2019-12-08 06:57:50,971 [main] WARN org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - Encountered Warning
2019-12-08 06:57:50,971 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - Success!
2019-12-08 06:57:50,987 [main] INFO org.apache.pig.Main - Pig script completed in 9 minutes, 26 seconds and 73 milliseconds (566073 ms)
```

MapReduce Runtime:

```
File Input Format Counters
  Bytes Read=4192925770
File Output Format Counters
  Bytes Written=0
Time taken for MapReduce Job =1816.814599371 seconds
```

Theoretically and practically Pig is supposed to top MapReduce in terms of processing time and it is quite evident Pig is faster than MapReduce and twice as simple. MapReduce consume approx. 1800 seconds ~ 30 mins as oppose to Pig's 9 minutes

Output:

AMERICAN DISTILLED SPIRIT SPECIALTY	881210.5191912651	55262	40359.770000100136
MISC. IMPORTED CORDIALS & LIQUEURS	3.7832800452870846E7	1877575	1453209.4611447453
AMERICAN CORDIALS & LIQUEURS	101481.1705095768	7657	4886.100000582635
STRAIGHT BOURBON WHISKIES	1.3097322989052832E8	7471154	6811651.311300226
AMERICAN FLAVORED VODKA	2.904975475186181E7	2870579	2116881.8510091044
WHITE CREME DE CACAO	241671.8572602272	35292	26446.5
IOWA DISTILLERIES	557729.2918243408	16185	12138.75
IMPORTED DRY GINS	3.1782569333756447E7	1374151	1318844.620017603
AMERICAN BRANDIES	1.6655681614335299E7	2528667	1362185.6645742469
SCOTCH WHISKIES	4.1172486617925644E7	1836787	2053085.460024152
LOW PROOF VODKA	142098.76273345947	8608	6966.0
VODKA 80 PROOF	1.45764888205423E8	18124615	1.779739678086917E7
GRAPE SCHNAPPS	1537955.1531925201	139783	133591.0
FLAVORED RUM	4.8862814052962065E7	3700204	3280113.430011064
WHITE RUM	2.088090985726595E7	1850419	1989563.5202885456
PUERTO RICO & VIRGIN ISLANDS RUM	3.971689315653813E7	3615422	3889666.9204342887
IMPORTED CORDIALS & LIQUEUR	71679.6097176075	3503	2473.0700000599027
AMERICAN CORDIALS & LIQUEUR	1.1005496307138085E7	2743191	835363.2027280573
HIGH PROOF BEER - AMERICAN	6237.839935302734	54	40.5
IOWA DISTILLERY WHISKIES	216305.1597366333	5799	4348.120000004768
TROPICAL FRUIT SCHNAPPS	430689.1462507248	60062	54243.25
NEUTRAL GRAIN SPIRITS	1912765.9047842026	144001	108154.95000010729
STRAWBERRY SCHNAPPS	543700.2474107742	72278	54208.5
SPECIAL ORDER ITEMS	5416297.283221245	322556	434679.76002436876
SCHNAPPS - IMPORTED	13483.799942016602	620	288.06000232696533
AMERICAN COCKTAILS	2.2720367726840496E7	2120834	3206271.1814655513
CINNAMON SCHNAPPS	1971718.8432440758	179177	151383.7699996233
CANADIAN WHISKIES	2.817220216301589E8	20082925	2.0877693360410452E7
AMERICAN DRY GINS	2.773394654225993E7	3889976	3039616.6554565094
WHISKEY LIQUEUR	1.0753896482602584E8	9189490	6346775.24666366
IRISH WHISKIES	4.057679356723404E7	1614906	1390139.9403407238
DELISTED ITEMS	4589.919988155365	149	111.98000001907349
CREAM LIQUEURS	5.4876392762699604E7	2998603	2515505.831994295
ANISETTE	37839.72998046875	5430	4072.5
MEZCAL	351645.24660873413	12743	9553.350000023842
IMPORTED DISTILLED SPIRITS SPECIALTY	19883.110072135925	568	454.5
TEMPORARY & SPECIALTY PACKAGES	1.5539515988339424E7	761342	631382.1708563119
NEUTRAL GRAIN SPIRITS FLAVORED	810302.4198875427	48912	34491.20002491772
MISCELLANEOUS SCHNAPPS	1775350.1104069948	179899	133371.15015863627
IMPORTED VODKA - MISC	2.963593110564089E7	1772283	1301889.440944627
BUTTERSCOTCH SCHNAPPS	2138725.056793213	221868	190362.5799999237
PEPPERMINT SCHNAPPS	6163636.847704172	878339	789853.1800096035
BLACKBERRY BRANDIES	3713928.2971582413	425294	373640.7300006151
AMARETTO - IMPORTED	70470.3021774292	1509	1131.75
BLENDED WHISKIES	6.425061144902766E7	6864961	6518767.41445785
APRICOT BRANDIES	1919003.7807621956	256932	180888.72001555562
IMPORTED VODKAS	5.4613833506902695E7	2926863	2909450.671165511
VODKA FLAVORED	3.870456769447374E7	3693136	2907612.7604818046
PEACH SCHNAPPS	5705569.960195065	567257	518755.75000166893
IMPORTED VODKA	7.751248114882612E7	3821869	3716375.7120935693
MIXTO TEQUILA	2.6883998862163544E7	2063063	1730605.341156412
FLAVORED GIN	346954.893491745	31282	23436.300000041723
BARBADOS RUM	854201.7081031799	59188	44283.0700000003
TRIPLE SEC	5616632.897033453	1409185	1372249.1599998474

4. Exploring expansion opportunities using Zip code analysis

The motivation behind this analysis is to provide a comprehensive overview of feasibility of an expansion. If someone is interested in opening a liquor store, one would naturally want to know if he/she is going to sustain and make profit or is there plenty of competition around.

Pig Script:

```
sales = LOAD 'hdfs://localhost:9000/project_data/Iowa_Liquor_Sales.csv'
        using org.apache.pig.piggybank.storage.CSVExcelStorage(',', 'NO_MULTILINE', 'NOCHANGE',
        'SKIP_INPUT_HEADER')
        as (Invoice:chararray, Date:chararray, Store_name:chararray, Store_number:int,
        Address:chararray, City:chararray, Zipcode:chararray, Store_location:chararray, County_Number:int,
        County:chararray, Category:int, Category_name:chararray, Vendor_number:int, Vendor_name:chararray,
        Item_number:int, Item_description:chararray, Pack:int, Bottle_volume:int, State_bottle_cost:float,
        State_bottle_retail:float, Bottles_sold:int, Sale:float, Volume_liters:float, Volume_gallons:float);

cleanData = FILTER sales BY Zipcode matches '5.*';

A = FOREACH cleanData GENERATE Zipcode, Store_number, Sale, Volume_liters;

uniq = DISTINCT A;

grpD = GROUP uniq BY Zipcode;

op = FOREACH grpD GENERATE FLATTEN(group),
        COUNT_STAR(uniq.Store_number) as stores,
        SUM(uniq.Sale) as total_sales,
        (SUM(uniq.Sale)/SUM(uniq.Volume_liters)) as avg_cost_per_litre;

mergedOp = FOREACH (GROUP op all) GENERATE FLATTEN(op);

STORE mergedOp INTO 'hdfs://localhost:9000/pig_outputs/zipcodeAnalysis_op';
```

Here, I encountered a couple of zip codes namely, 80904: Colorado & 87325: New Mexico. In order to skip such dirty data, we perform filtering before we group data by Zip code.

Additionally, we want to get a count of stores within a zip code. As our data is on a transactional level, simply taking a count can result in incorrect values. We want distinct number of stores within a zip code where transactions have taken place. For this, we get a distinct of zip codes.

Since, some rows have nulls in place of store number, we take a COUNT_STAR that considers nulls but does not include it in the count.

PIG Runtime:

```
2019-12-11 19:53:22,911 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.Ma
pReduceLauncher - Success!
2019-12-11 19:53:22,944 [main] INFO org.apache.pig.Main - Pig script completed in 7 minutes, 59 sec
onds and 442 milliseconds (479442 ms)
karmickoala@appliedore:~$
```

Once again, Pig has topped its previous best with 7 minutes this time.

Output:

57222	124	12075.060056209564	20.465509615105567
56201	215	40544.37019729614	16.426362887590862
52807	28894	1.687030425382352E7	16.285295872906193
52806	18470	3306699.8369580507	14.746714276878066
52804	30926	7668303.421343803	17.494932509633543
52803	11327	1888296.3840950727	20.62481388806394
52802	16050	1620849.8925184011	16.946415279588358
52801	1250	127026.26051592827	22.303856205433398
52778	3475	149756.9301867485	13.771967927940755
52777	446	28895.040053844452	14.972609423512745
52776	3898	310172.0903711319	18.713824090241538
52773	91	11279.880031585693	16.533352920641256
52772	3889	497625.65063619614	13.04117088049011
52768	474	31113.22997045517	14.009225913633966
52767	52	4136.180023193359	15.130336258994083
52766	82	9981.040019989014	17.015069928339003
52761	31338	5371147.442671418	15.941605049459257
52755	92	11267.040031433105	18.59186170764126
52753	6554	1055497.7405002117	22.266170215019166
52751	1316	109504.28019475937	14.901359575240557
52748	8054	952018.0905399323	15.523609210890614
52747	3051	107689.83007013798	12.76422243582715
52742	3235	1015584.8493480682	12.841542900492144
52738	1977	117835.9401705265	17.13050391747908
52733	2828	461002.79074668884	14.901986585184694
52732	19908	3445747.757604003	15.873690695554757
52730	1533	132394.97010695934	15.68978666672715
52728	474	57914.76009893417	16.132245134157436
52726	2809	102385.35000896454	14.29935028584452
52722	25053	6680384.384613633	17.004001163858128
52671	578	31905.960065841675	16.76834061040624
52659	231	13147.200013875961	14.58094424825562
52656	1420	191029.38041996956	15.679351616528056
52655	5012	912816.5121921301	16.61253032313919
52653	3201	194155.72993278503	14.553933403592948
52645	213	18646.280005455017	16.9978304158356
52641	7629	1798594.1522603035	15.25862936198582
52639	725	43902.11004972458	15.063290226416582
52638	87	11167.140022277832	17.392944509104417
52637	2963	149788.3299444914	14.261629883757085
52632	13033	3616609.6379024982	16.609233648886864
52627	13118	2975090.1994885206	15.46119206998359
52626	508	59579.68987989426	14.956430777474589
52625	371	50465.82000350952	16.294119150061537
52623	269	23223.030084609985	17.019692511154826
52601	25547	3738041.055680752	16.6608934449227
52591	4457	231626.54983568192	13.386047844362912
52590	175	18953.100069999695	14.717539376542222
52577	15923	2016656.402243495	14.916743075687123
52571	1272	108700.05999183655	18.89532136552108
52565	2479	209526.39037299156	15.917720702485635
52561	79	9647.15002822876	17.26961087706565
52556	8799	1715332.5036969185	17.219810721526013
52554	424	35645.440104961395	14.772189131448789
-----	----	-----	-----

5. Researching top alcohol brand sold by each store in the state

The intent behind this use case was to understand the popularity of liquor by stores across the states. Any store near a college, university or any institution that has young students around would sell a lot of beers or alcopops. Similarly, stores right in middle of a community will cater to a larger older population and may sell hard liquors as opposed to beers.

This supplements well to the 2nd use case where we found out top 5 brands per year based on sales. It might also help in stocking inventory.

Hive Script:

```
INSERT OVERWRITE DIRECTORY 'hdfs://localhost:9000/hive_outputs/MaxLiquorPerStore_op' ROW
FORMAT DELIMITED FIELDS TERMINATED BY '\t'
|
SELECT t1.Store_name, t1.Item_description
FROM(
    SELECT Store_name, Item_description, SUM(Bottles_sold) AS totalSold
    FROM sales
    GROUP BY Item_description, Store_name
) AS t1
JOIN(
    SELECT s.Store_name, MAX(s.totalSold) AS maxSold
    FROM(
        SELECT Store_name, Item_description, SUM(
            Bottles_sold) AS totalSold
        FROM sales
        GROUP BY Item_description, Store_name
        ) s
    GROUP BY s.Store_name
)AS t2
ON t1.totalSold = t2.maxSold AND
t1.store name = t2.store name
GROUP BY t1.Store_name, t1.Item_description
```

The SQL here is self-explanatory.

1. First, we get the store name, item description and total sum of bottles sold which is our metric
2. Then, we get the value maximum bottles have been sold in a store
3. Join both queries and where ever the maximum count matches the sum of bottles sold for an alcohol, that's our most popular liquor for that store

Hive Run time:

```
Stage-Stage-4: Map: 1 Cumulative CPU: 9.1 sec HDFS Read: 71782155 HDFS Write: 105423 SUCCESS
Total MapReduce CPU Time Spent: 26 minutes 42 seconds 740 msec
OK
Time taken: 712.361 seconds
hive> █
```

Hive seems to be pretty fast as well. 712 seconds ~ 12 minutes to crunch almost 17 million records. Impressive!

Output:

Casey's #3746	Hawkeye Vodka	
Casey's General Store # 1591/	Decorah	Admiral Nelson Spiced Rum
Casey's General Store # 2417/	Newton	Black Velvet
Casey's General Store # 2618/	Fredricksburg	Black Velvet
Casey's General Store # 2653 /	Toledo	Black Velvet
Casey's General Store # 3518/	Des Moines	Mccormick Vodka Pet
Casey's General Store # 3546/	Monona	Mccormick Vodka
Casey's General Store #1020 /	Le Cl	64858
Casey's General Store #1374 /	Colfax	Hawkeye Vodka
Casey's General Store #1417 /	Iowa Falls	Smirnoff 80prf
Casey's General Store #2060 /	Pocahontas	Black Velvet
Casey's General Store #2168	Black Velvet	
Casey's General Store #2168	Hawkeye Vodka	
Casey's General Store #2304 /	Slater	Mccormick Vodka Pet
Casey's General Store #2319 /	Ft Madison	Hawkeye Vodka
Casey's General Store #2488 /	N English	Black Velvet
Casey's General Store #2521 /	Adair	Mccormick Vodka
Casey's General Store #2523 /	Monroe	Hawkeye Vodka
Casey's General Store #2526 /	Wellsb	Five O'clock Vodka
Casey's General Store #2550 /	Osceola	Mccormick Vodka Pet
Casey's General Store #2644 /	Earlha	Hawkeye Vodka
Casey's General Store #2782 /	Cedar Rapids	Hawkeye Vodka
Casey's General Store #2813 /	Fort Dodge	Five O'clock Vodka
Casey's General Store #2816 /	Johnston	Black Velvet
Casey's General Store #2824 /	WDM	Mccormick Vodka Pet
Casey's General Store #2902 /	Spence	Five O'clock PET Vodka
Casey's General Store #3024 /	Mediapolis	Black Velvet
Casey's General Store #3050 /	Counci	Inc."
Casey's General Store #3210 /	Urband	Burnett's Vodka 80 Prf
Casey's General Store #3291	Black Velvet	
Casey's General Store #3333 /	Pleasant Hill	Hawkeye Vodka
Casey's General Store #3606	Hawkeye Vodka	
Casey's General Store #3730 /	Paullina	Black Velvet
Casey's General Store #95 /	Dexter	Black Velvet
Circle K #4706604 /	Burlington	Five O'Clock Vodka
Cork 'N Bottle /	Manchester	Captain Morgan Spiced Rum
County Market #214 /	Fort Madison	Hawkeye Vodka
Creekside Market	Hawkeye Vodka	
Crossroads of Algona	Five O'clock Vodka	
Discount Tobacco and More /	Davenpor	Mccormick Vodka
Dyersville Liquor Mart	43338	
FRANKLIN STREET FLORAL & GIFT	4032	
Family Fare #791	Barton Vodka	
Fareway Stores #067 /	Evansdale	Five O'clock Vodka
Fareway Stores #657 /	Indianola	Hawkeye Vodka
Fareway Stores #788 /	Spencer	Black Velvet
Fareway Stores #850 /	Spirit Lake	Black Velvet
Fareway Stores #922 /	New Hampton	Black Velvet
Fareway Stores #989	Titos Handmade	Vodka
Fill R Up	Fireball Cinnamon	
Food Land Super Markets /	Missouri V	4356
Foodland Super Markets /	Woodbine	Hawkeye Vodka
Gasland / Burlington	Five O'clock Vodka	
Gasland N8th St / Burlington	Five O'Clock Vodka	